

**Lab 12 Report**  
**Nguyễn Khánh Nam**  
**20225749**

<b>Code: row-major.asm</b> .....	<b>1</b>
<b>Code: colum-major</b> .....	<b>2</b>
<b>Result:</b> .....	<b>4</b>
row-major:.....	4
column-major.....	4

Code: row-major.asm

#####  
#####

# Row-major order traversal of 16 x 16 array of words.

# Pete Sanderson

# 31 March 2007

#

# To easily observe the row-oriented order, run the Memory Reference

# Visualization tool with its default settings over this program.

# You may, at the same time or separately, run the Data Cache Simulator

# over this program to observe caching performance. Compare the results

# with those of the column-major order traversal algorithm.

#

# The C/C++/Java-like equivalent of this MIPS program is:

#   int size = 16;

#   int[size][size] data;

#   int value = 0;

#   for (int row = 0; col < size; row++) {

#     for (int col = 0; col < size; col++) {

#       data[row][col] = value;

#       value++;

#     }

# }  
#

# Note: Program is hard-wired for 16 x 16 matrix. If you want to change this,

#   three statements need to be changed.

#   1. The array storage size declaration at "data:" needs to be changed from

#     256 (which is 16 \* 16) to #columns \* #rows.

#   2. The "li" to initialize \$t0 needs to be changed to new #rows.

#   3. The "li" to initialize \$t1 needs to be changed to new #columns.

```

#
.data
data: .word 0 : 256 # storage for 16x16 matrix of words
.text
li $t0, 16 # $t0 = number of rows
li $t1, 16 # $t1 = number of columns
move $s0, $zero # $s0 = row counter
move $s1, $zero # $s1 = column counter
move $t2, $zero # $t2 = the value to be stored
# Each loop iteration will store incremented $t1 value into next element of matrix.
# Offset is calculated at each iteration. offset = 4 * (row*#cols+col)
# Note: no attempt is made to optimize runtime performance!
loop: mult $s0, $t1 # $s2 = row * #cols (two-instruction sequence)
mflo $s2 # move multiply result from lo register to $s2
add $s2, $s2, $s1 # $s2 += column counter
sll $s2, $s2, 2 # $s2 *= 4 (shift left 2 bits) for byte offset
sw $t2, data($s2) # store the value in matrix element
addi $t2, $t2, 1 # increment value to be stored
# Loop control: If we increment past last column, reset column counter and
increment row counter
# If we increment past last row, we're finished.
addi $s1, $s1, 1 # increment column counter
bne $s1, $t1, loop # not at end of row so loop back
move $s1, $zero # reset column counter
addi $s0, $s0, 1 # increment row counter
bne $s0, $t0, loop # not at end of matrix so loop back
# We're finished traversing the matrix.
li $v0, 10 # system service 10 is exit
syscall # we are outta here.

```

Code: colum-major

```

#####
#
# Column-major order traversal of 16 x 16 array of words.
# Pete Sanderson
# 31 March 2007
#
# To easily observe the column-oriented order, run the Memory Reference
# Visualization tool with its default settings over this program.
# You may, at the same time or separately, run the Data Cache Simulator
# over this program to observe caching performance. Compare the results
# with those of the row-major order traversal algorithm.
#

```

# The C/C++/Java-like equivalent of this MIPS program is:

```
# int size = 16;
# int[size][size] data;
# int value = 0;
# for (int col = 0; col < size; col++) {
#     for (int row = 0; row < size; row++) {
#         data[row][col] = value;
#         value++;
#     }
# }
```

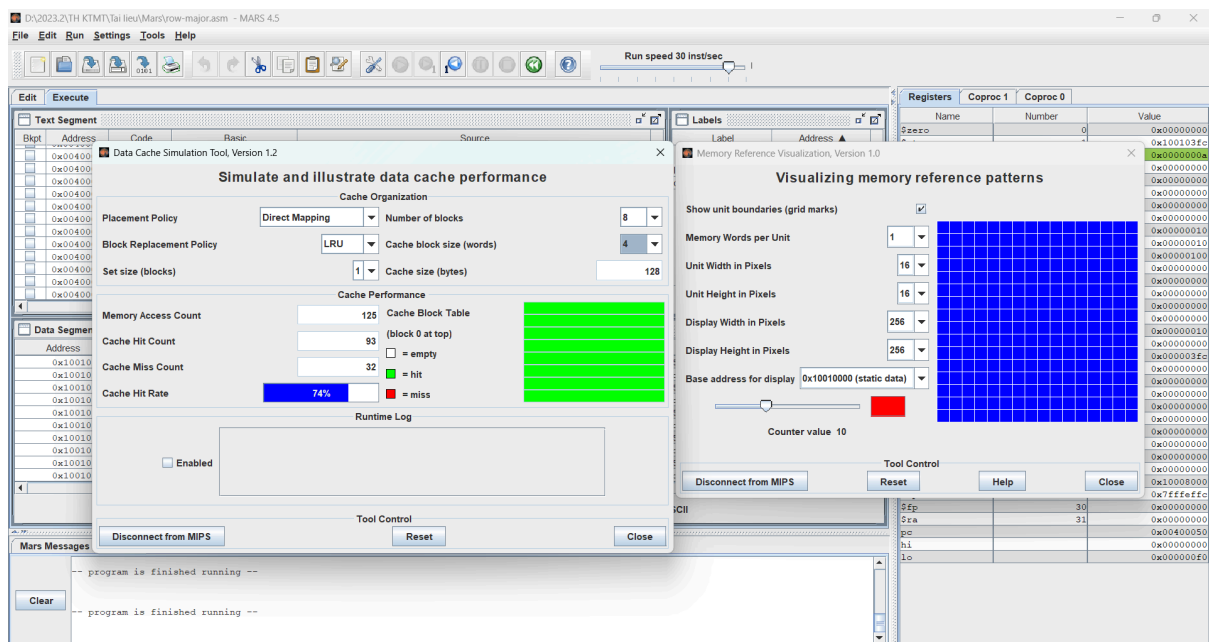
# Note: Program is hard-wired for 16 x 16 matrix. If you want to change this,  
# three statements need to be changed.  
# 1. The array storage size declaration at "data:" needs to be changed from  
# 256 (which is 16 \* 16) to #columns \* #rows.  
# 2. The "li" to initialize \$t0 needs to be changed to the new #rows.  
# 3. The "li" to initialize \$t1 needs to be changed to the new #columns.  
#

```
.data
data: .word 0 : 256 # 16x16 matrix of words
.text
li $t0, 16 # $t0 = number of rows
li $t1, 16 # $t1 = number of columns
move $s0, $zero # $s0 = row counter
move $s1, $zero # $s1 = column counter
move $t2, $zero # $t2 = the value to be stored
# Each loop iteration will store incremented $t1 value into next element of matrix.
# Offset is calculated at each iteration. offset = 4 * (row*#cols+col)
# Note: no attempt is made to optimize runtime performance!
loop: mult $s0, $t1 # $s2 = row * #cols (two-instruction sequence)
mflo $s2 # move multiply result from lo register to $s2
add $s2, $s2, $s1 # $s2 += col counter
sll $s2, $s2, 2 # $s2 *= 4 (shift left 2 bits) for byte offset
sw $t2, data($s2) # store the value in matrix element
addi $t2, $t2, 1 # increment value to be stored
# Loop control: If we increment past bottom of column, reset row and increment
column
# If we increment past the last column, we're finished.
addi $s0, $s0, 1 # increment row counter
bne $s0, $t0, loop # not at bottom of column so loop back
move $s0, $zero # reset row counter
addi $s1, $s1, 1 # increment column counter
bne $s1, $t1, loop # loop back if not at end of matrix (past the last column)
# We're finished traversing the matrix.
```

```
li    $v0, 10    # system service 10 is exit
syscall          # we are outta here.
```

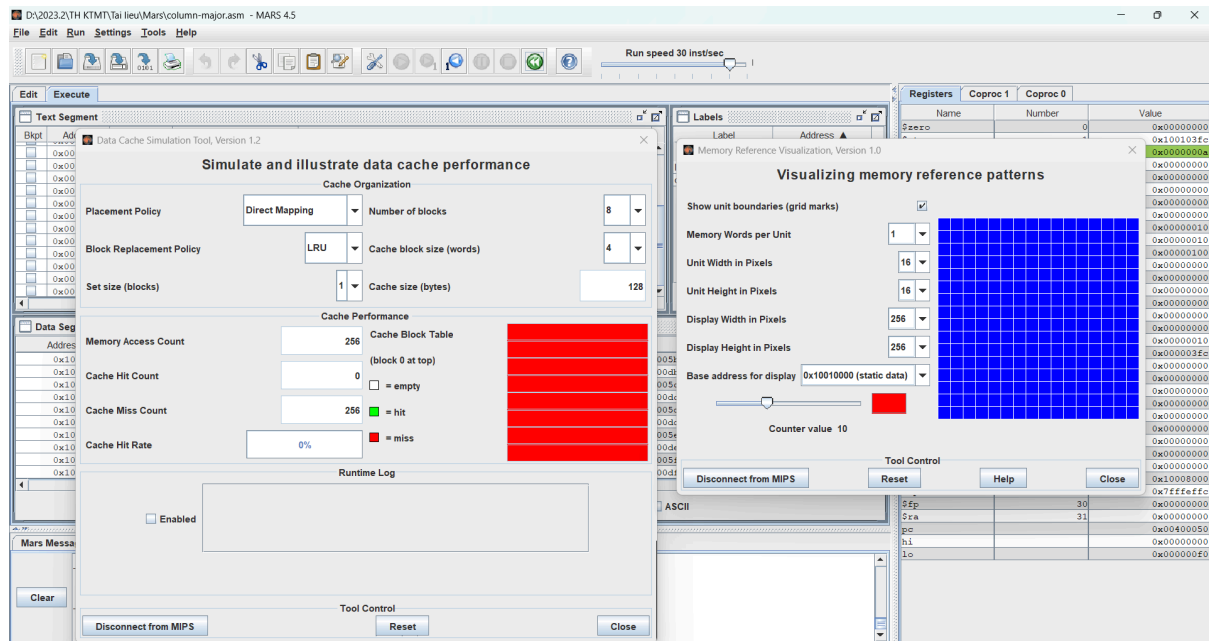
Result:

row-major:



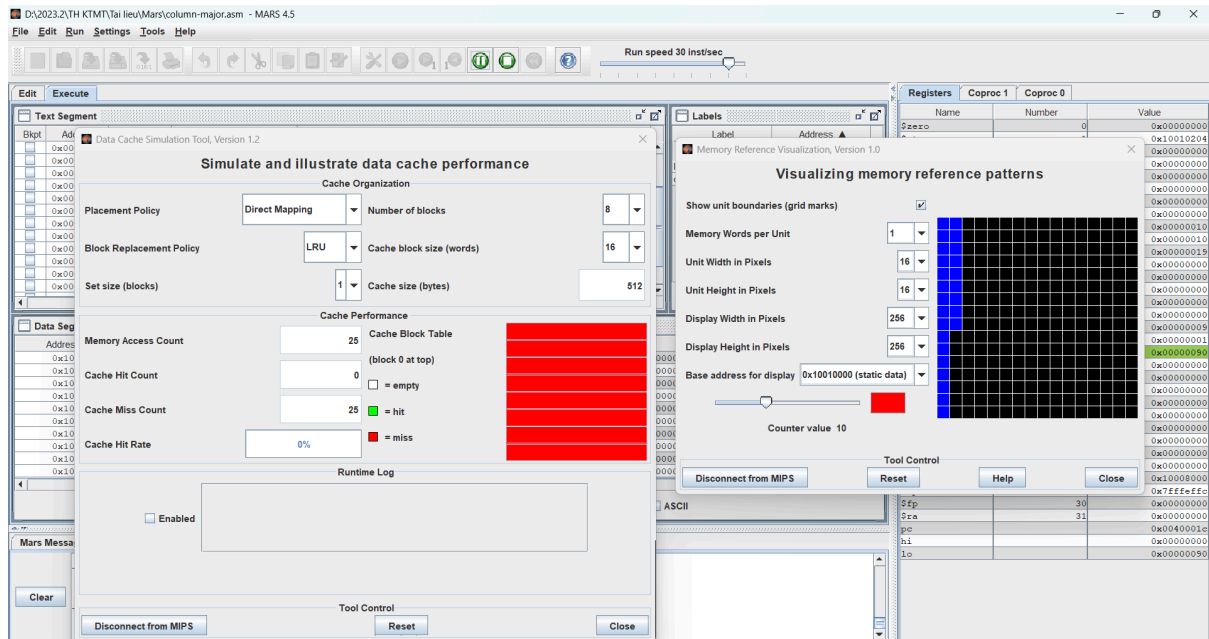
- The final cache hit rate:  $\frac{3}{4}$ 
  - + Mỗi lần cache miss, sẽ có 1 khối 4 words được đọc lưu vào cache.
  - + Trong truy vấn theo hàng, các phần tử mảng cũng được truy cập lần lượt theo đúng thứ tự lưu nên trong 1 khối có 4 word thì sẽ có 3 words trùng đã được đọc trong khối từ trước đó trong lần truy cập vào ô tiếp theo.
  - + Sau đó đến ô thứ 5 -> cache miss nên tự lưu trước 4 ô từ ô thứ 5.
  - + Lặp lại như trên.
- Increase block size to 8 words: cache hits =  $\frac{7}{8}$
- Decrease to 2 words: cache hits =  $\frac{1}{2}$

column-major

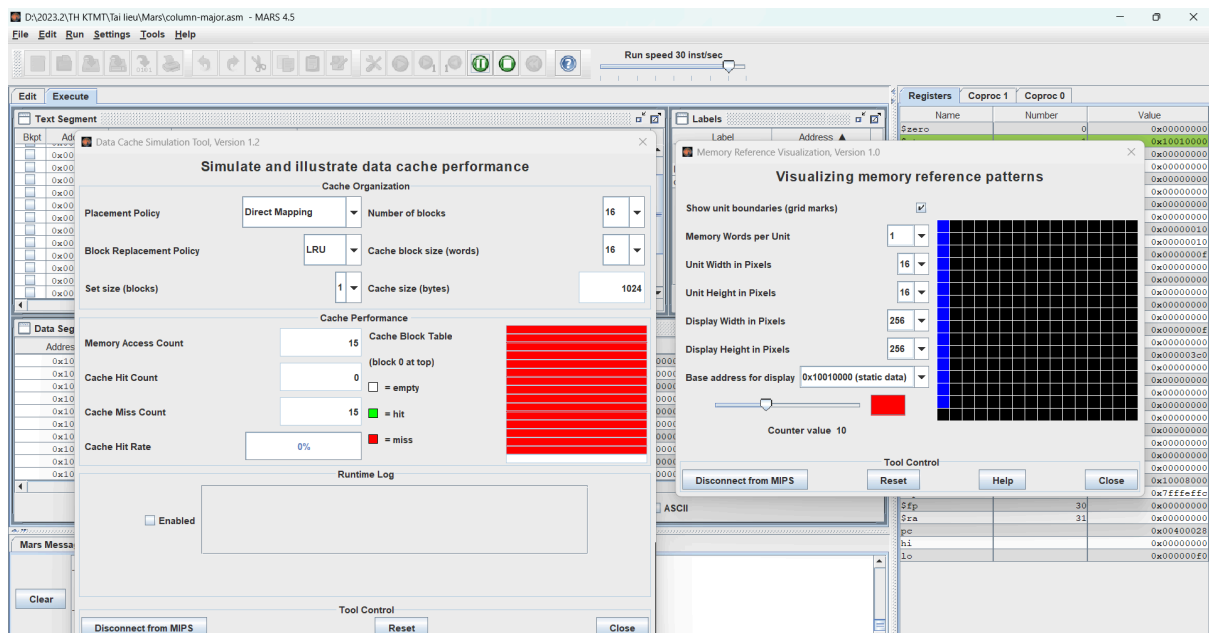


- The final cache hit rate: 0/4
  - + Mỗi lần cache miss, sẽ có 1 khối 4 words được đọc lưu vào cache.
  - + Do truy cập dữ liệu theo cột nên sẽ liên tục bị cache miss do 1 khối 4 words được lưu theo hàng mỗi lần cache miss.
- Increase block size to 8 words: cache hits = 0
- Decrease to 2 words: cache hits = 0
- The cache performance for this program: Truy cập từng words theo cột nên mỗi word được truy cập sẽ cách word trong block được lưu vào cache 16 words -> cache miss (dữ liệu trong block bị thay thế dù chưa được truy cập)

- Block size tăng lên 16 nhưng số lượng block vẫn là 8 nên không đủ để cho dữ liệu truy cập được cache hits. Do chưa kịp truy cập tới dữ liệu cần dùng thì đã bị thay thế sau khi hết 8 block



- Block size tăng lên 16 cùng số lượng Block cũng được tăng lên 16 mà mảng 16x16 nên các block sẽ được lưu hết vào cache có đủ số lượng block 16 words cho cả mảng nên sau khi duyệt hết cột đầu tiên từ cột thứ 2 sẽ có cache hits lần lượt các block.



D:\2023\2\TH KTM\Tai lieu\Mars\column-major.asm - MARS 4.5

File Edit Run Settings Tools Help

Run speed 30 inst/sec

Edit Execute

Text Segment

Data Cache Simulation Tool, Version 1.2

### Simulate and illustrate data cache performance

Cache Organization

Placement Policy: Direct Mapping Number of blocks: 16

Block Replacement Policy: LRU Cache block size (words): 16

Set size (blocks): 1 Cache size (bytes): 1024

Cache Performance

Memory Access Count: 27 Cache Block Table (block 0 at top)

Cache Hit Count: 11 = empty

Cache Miss Count: 16 = hit

Cache Hit Rate: 41%

Runtime Log

Enabled

Tool Control

Disconnect from MIPS Reset Close

Memory Reference Visualization, Version 1.0

### Visualizing memory reference patterns

Show unit boundaries (grid marks): ☒

Memory Words per Unit: 1

Unit Width in Pixels: 16

Unit Height in Pixels: 16

Display Width in Pixels: 256

Display Height in Pixels: 256

Base address for display: 0x10010000 (static data)

Counter value 10

Tool Control

Disconnect from MIPS Reset Help Close

Registers

Name	Number	Value
\$zero	0	0x00000000
\$at	1	0x10010284
\$v0	2	0x00000000
\$v1	3	0x00000000
\$a0	4	0x00000000
\$a1	5	0x00000000
\$a2	6	0x00000000
\$a3	7	0x00000000
\$t0	8	0x00000000
\$t1	9	0x00000000
\$t2	10	0x00000000
\$t3	11	0x00000000
\$t4	12	0x00000000
\$t5	13	0x00000000
\$t6	14	0x00000000
\$t7	15	0x00000000
\$s0	16	0x00000000
\$s1	17	0x00000000
\$s2	18	0x00000000
\$s3	19	0x00000000
\$s4	20	0x00000000
\$s5	21	0x00000000
\$s6	22	0x00000000
\$s7	23	0x00000000
\$ra	24	0x00000000
\$gp	25	0x00000000
\$sp	26	0x00000000
\$fp	27	0x00000000
\$hi	28	0x00000000
\$lo	29	0x00000000

D:\2023\2\TH KTM\Tai lieu\Mars\column-major.asm - MARS 4.5

File Edit Run Settings Tools Help

Run speed 30 inst/sec

Edit Execute

Text Segment

Data Cache Simulation Tool, Version 1.2

### Simulate and illustrate data cache performance

Cache Organization

Placement Policy: Direct Mapping Number of blocks: 16

Block Replacement Policy: LRU Cache block size (words): 16

Set size (blocks): 1 Cache size (bytes): 1024

Cache Performance

Memory Access Count: 256 Cache Block Table (block 0 at top)

Cache Hit Count: 240 = empty

Cache Miss Count: 16 = hit

Cache Hit Rate: 94%

Runtime Log

Enabled

Tool Control

Disconnect from MIPS Reset Close

Memory Reference Visualization, Version 1.0

### Visualizing memory reference patterns

Show unit boundaries (grid marks): ☒

Memory Words per Unit: 1

Unit Width in Pixels: 16

Unit Height in Pixels: 16

Display Width in Pixels: 256

Display Height in Pixels: 256

Base address for display: 0x10010000 (static data)

Counter value 10

Tool Control

Disconnect from MIPS Reset Help Close

Registers

Name	Number	Value
\$zero	0	0x00000000
\$at	1	0x100103fc
\$v0	2	0x00000000
\$v1	3	0x00000000
\$a0	4	0x00000000
\$a1	5	0x00000000
\$a2	6	0x00000000
\$a3	7	0x00000000
\$t0	8	0x00000000
\$t1	9	0x00000000
\$t2	10	0x00000000
\$t3	11	0x00000000
\$t4	12	0x00000000
\$t5	13	0x00000000
\$t6	14	0x00000000
\$t7	15	0x00000000
\$s0	16	0x00000000
\$s1	17	0x00000000
\$s2	18	0x00000000
\$s3	19	0x00000000
\$s4	20	0x00000000
\$s5	21	0x00000000
\$s6	22	0x00000000
\$s7	23	0x00000000
\$ra	24	0x00000000
\$gp	25	0x00000000
\$sp	26	0x00000000
\$fp	27	0x00000000
\$hi	28	0x00000000
\$lo	29	0x00000000